


# Exact Inference

Francisco Pereira

Filipe Rodrigues



# Outline

- Recap
- Analytical Derivation (Bayes' theorem)
- Variable Elimination
- Belief Propagation

## Previously, in MBML...


- Representation (weeks 1-4)
- Modelling toolbox (weeks 5-8, +12)
- Inference (weeks 9-11)

# Previously, in MBML...

- Representation (weeks 1-4)
  - PGM basics
  - Conditional independence
  - Generative processes
  - Joint probability distribution
  - Priors and likelihood
  - Factorization

$$p(\beta, \mathbf{z}, \mathbf{k}, \mathbf{dt}) = p(\beta | \sigma_\beta) \prod_{n=1}^N p(k_n | \mathbf{x}_n, \beta, \sigma_\epsilon) p(z_n | \pi) p(dt_n | z_n, k_n)$$

- 1 Draw a pair of parameters<sup>1</sup>,  $\beta \sim \mathcal{N}(\mathbf{0}, I\sigma_\beta)$
- 2 For  $n = 1..N$ 
  - 1 Draw one value for  $z_n$ , such that  $z_n \sim \text{Bern}(\pi)$ .
    - If  $z_n = 1$ , the bus has stopped ( $z_n = 0$  otherwise)
    - Distributed as Bernoulli, with parameter  $\pi$
  - 2 Draw one value for  $k_n$ , such that  $k_n \sim \mathcal{N}(\mathbf{x}_n^T \beta, \sigma_\epsilon)$
  - 3 If  $z_n = 1$ ,  $dt_n = k_n$ ,
    - otherwise  $dt_n = 0$



## Previously, in MBML...

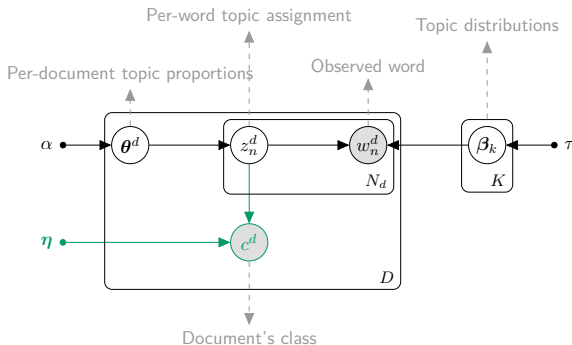
- Representation (weeks 1-4)
- Modelling toolbox (weeks 5-8, +12)
  - Mixture models
  - Different likelihoods (Gaussian, Poisson, etc.)
  - Link functions (log, softmax, Probit, etc.)
  - Non-linear relationships (e.g. using neural networks)
  - Discrete vs. continuous target variables
  - Heteroscedastic models
  - Hierarchical models
  - Temporal models (continuous and discrete)
  - Topic modelling (discrete data; e.g. text corpora)
  - Bayesian non-parametric models (week 12)
- Inference (weeks 9-11)

## Previously, in MBML...

- Representation (weeks 1-4)
- Modelling toolbox (weeks 5-8, +12)
  - Bayesian Gaussian Mixture models
  - Bayesian Linear regression
  - Poisson regression
  - Heteroscedastic regression
  - Bayesian Logistic regression
  - Bayesian Probit regression
  - Hierarchical Logistic regression
  - Autoregressive models
  - Linear dynamical systems (e.g. Kalman filter)
  - Hidden Markov models
  - Latent Dirichlet allocation
  - Gaussian processes (week 12)
- Inference (weeks 9-11)

## Previously, in MBML...

- Representation (weeks 1-4)
- Modelling toolbox (weeks 5-8, +12)
  - Mix & Match (e.g. LDA + (Logistic) Regression = Supervised LDA)




$$p(c^d | \bar{z}, \eta) = \frac{\exp(\eta_c^T \bar{z})}{\sum_{l=1}^C \exp(\eta_l^T \bar{z})}$$

- Inference (weeks 9-11)

## Previously, in MBML...

- Representation (weeks 1-4)
- Modelling toolbox (weeks 5-8, +12)
- Inference (weeks 9-11)
  - Exact inference
    - Analytical (Bayes' rule)
    - Variable elimination
    - Belief propagation
  - Approximate inference
    - Stochastic methods - Markov chain Monte Carlo (MCMC)
    - Deterministic methods - Variational inference (VI)





# Exact inference

- Computation of the exact posterior probability distribution over the variables of interest
  - Best possible solution, given data and specification
- Analytical derivations:
  - High computational efficiency
  - However, quite often, not possible at all  
(when it is not possible, we say that it is *intractable*)
- Algorithmic methods
  - Variable elimination
  - Message passing

## Analytical derivation

- The Gaussian form has very nice properties<sup>1</sup>
- A multivariate Gaussian with  $d$  dimensions is:

$$\mathcal{N}(\mathbf{x}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{(2\pi)^{D/2}|\boldsymbol{\Sigma}|^{1/2}} e^{-\frac{1}{2}(\mathbf{x}-\boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1}(\mathbf{x}-\boldsymbol{\mu})}$$

- The inverse of the covariance matrix is  $\boldsymbol{\Lambda} = \boldsymbol{\Sigma}^{-1}$  is called *precision matrix*
- If you have a marginal Gaussian distribution for  $\mathbf{x}$  and a conditional Gaussian distribution of  $\mathbf{y}$  given  $\mathbf{x}$  in the form

$$p(\mathbf{x}) = \mathcal{N}(\mathbf{x}|\boldsymbol{\mu}, \boldsymbol{\Lambda}^{-1})$$

$$p(\mathbf{y}|\mathbf{x}) = \mathcal{N}(\mathbf{y}|\mathbf{Ax} + \mathbf{b}, \mathbf{L}^{-1})$$

- Then the marginal distribution for  $\mathbf{y}$  and a conditional Gaussian distribution of  $\mathbf{x}$  given  $\mathbf{y}$  have the form

$$p(\mathbf{y}) = \mathcal{N}(\mathbf{y}|\mathbf{A}\boldsymbol{\mu} + \mathbf{b}, \mathbf{L}^{-1} + \mathbf{A}\boldsymbol{\Lambda}^{-1}\mathbf{A}^T)$$

$$p(\mathbf{x}|\mathbf{y}) = \mathcal{N}(\mathbf{x}|\boldsymbol{\Gamma}[\mathbf{A}^T\mathbf{L}(\mathbf{y} - \mathbf{b})] + \boldsymbol{\Gamma}\boldsymbol{\mu}, \boldsymbol{\Gamma})$$

- where  $\boldsymbol{\Gamma} = (\boldsymbol{\Lambda} + \mathbf{A}^T\mathbf{L}\mathbf{A})^{-1}$

---

<sup>1</sup>Check appendix of Bishop's book for many more useful properties!

## Analytical derivation

- If we have a joint Gaussian distribution  $N(\mathbf{x}|\boldsymbol{\mu}, \boldsymbol{\Sigma})$ , with  $\boldsymbol{\Lambda} = \boldsymbol{\Sigma}^{-1}$ , and we define the following partitions:

$$\mathbf{x} = \begin{pmatrix} \mathbf{x}_a \\ \mathbf{x}_b \end{pmatrix}, \boldsymbol{\mu} = \begin{pmatrix} \boldsymbol{\mu}_a \\ \boldsymbol{\mu}_b \end{pmatrix}$$
$$\boldsymbol{\Sigma} = \begin{pmatrix} \boldsymbol{\Sigma}_{aa} & \boldsymbol{\Sigma}_{ab} \\ \boldsymbol{\Sigma}_{ba} & \boldsymbol{\Sigma}_{bb} \end{pmatrix}, \boldsymbol{\Lambda} = \begin{pmatrix} \boldsymbol{\Lambda}_{aa} & \boldsymbol{\Lambda}_{ab} \\ \boldsymbol{\Lambda}_{ba} & \boldsymbol{\Lambda}_{bb} \end{pmatrix}$$

- Then the conditional distribution  $p(\mathbf{x}_a|\mathbf{x}_b)$  is given by

$$p(\mathbf{x}_a|\mathbf{x}_b) = \mathcal{N}(\mathbf{x}|\boldsymbol{\mu}_{a|b}, \boldsymbol{\Lambda}_{aa}^{-1})$$

$$\boldsymbol{\mu}_{a|b} = \boldsymbol{\mu}_a - \boldsymbol{\Lambda}_{aa}^{-1} \boldsymbol{\Lambda}_{ab}(\mathbf{x}_b - \boldsymbol{\mu}_b)$$

- And the marginal distribution  $p(\mathbf{x}_a)$  is given by

$$p(\mathbf{x}_a) = \mathcal{N}(\mathbf{x}_a|\boldsymbol{\mu}_a, \boldsymbol{\Sigma}_{aa})$$

## Analytical derivation

- If we have two (univariate) Gaussian distributions (for the **same** variable  $x$ ):

$$p(x) = \mathcal{N}(x|\mu_a, \sigma_a^2), \quad p(x) = \mathcal{N}(x|\mu_b, \sigma_b^2)$$

- Their product is:

$$p(x) = \mathcal{N}(x|\mu_{ab}, \sigma_{ab}^2)$$

- where

$$\mu_{ab} = \frac{\mu_a \sigma_b^2 + \mu_b \sigma_a^2}{\sigma_a^2 + \sigma_b^2}, \quad \sigma_{ab}^2 = \frac{\sigma_a^2 \sigma_b^2}{\sigma_a^2 + \sigma_b^2}$$

- This is very useful to directly combine models into a single prediction (e.g. ensemble models)!

## Analytical derivation

- Example of Bayesian linear regression
- The joint probability of our model is given by:

$$p(\mathbf{y}, \boldsymbol{\beta} | \mathbf{X}, \lambda, \sigma) = p(\boldsymbol{\beta} | \mathbf{0}, \lambda \mathbf{I}) \prod_{n=1}^N p(y_n | \boldsymbol{\beta}^T \mathbf{x}_n, \sigma^2)$$

- Since both our prior  $p(\boldsymbol{\beta} | \lambda)$  and likelihood  $p(y_n | \mathbf{x}_n, \boldsymbol{\beta}, \sigma)$  are both Gaussian, we apply Bayes' theorem to compute posterior over  $\boldsymbol{\beta}$ :

$$\begin{aligned} \underbrace{p(\boldsymbol{\beta} | \mathbf{y}, \mathbf{X}, \lambda, \sigma)}_{\text{posterior}} &\propto \underbrace{p(\boldsymbol{\beta} | \mathbf{0}, \lambda \mathbf{I})}_{\text{prior}} \underbrace{\prod_{n=1}^N p(y_n | \boldsymbol{\beta}^T \mathbf{x}_n, \sigma^2)}_{\text{likelihood}} \\ &= \mathcal{N}(\boldsymbol{\beta} | \mathbf{0}, \lambda \mathbf{I}) \prod_{n=1}^N \mathcal{N}(y_n | \boldsymbol{\beta}^T \mathbf{x}_n, \sigma^2) \\ &= \mathcal{N}(\boldsymbol{\beta} | \boldsymbol{\mu} = ?, \boldsymbol{\Sigma} = ?) \end{aligned}$$

- Can you do it?  $\boldsymbol{\mu} = ?$ ,  $\boldsymbol{\Sigma} = ?$  (Hint: look at slide 4...)

## Analytical derivation

$$\begin{aligned}
 \underbrace{p(\boldsymbol{\beta}|\mathbf{y}, \mathbf{X}, \lambda, \sigma)}_{\text{posterior}} &\propto \underbrace{p(\boldsymbol{\beta}|\mathbf{0}, \lambda\mathbf{I})}_{\text{prior}} \underbrace{\prod_{n=1}^N p(y_n|\boldsymbol{\beta}^T \mathbf{x}_n, \sigma^2)}_{\text{likelihood}} \\
 &= \mathcal{N}(\boldsymbol{\beta}|\mathbf{0}, \lambda\mathbf{I}) \prod_{n=1}^N \mathcal{N}(y_n|\boldsymbol{\beta}^T \mathbf{x}_n, \sigma^2)
 \end{aligned}$$

- Notice that, because our observations  $y_n$  are i.i.d., we can re-write:

$$\begin{aligned}
 &= \mathcal{N}(\boldsymbol{\beta}|\mathbf{0}, \lambda\mathbf{I}) \mathcal{N}(\mathbf{y}|\mathbf{X}\boldsymbol{\beta}, \sigma^2\mathbf{I}) \\
 &= \mathcal{N}(\boldsymbol{\beta}|\boldsymbol{\mu}=?, \boldsymbol{\Sigma}=?)
 \end{aligned}$$

- where  $\mathbf{y} = \{y_1, \dots, y_N\}$  and  $\mathbf{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$
- Thus, we can apply the properties in slide 4, yielding:

$$\begin{aligned}
 \boldsymbol{\mu} &= \boldsymbol{\Sigma} (\sigma^{-2} \mathbf{X}^T \mathbf{y}) \\
 \boldsymbol{\Sigma} &= (\lambda^{-1} \mathbf{I} + \sigma^{-2} \mathbf{X}^T \mathbf{X})^{-1}
 \end{aligned}$$

# Variable Elimination

## simple PGM

- A chain graph



$$p(z_1, z_2, z_3, z_4) = p(z_4|z_3) p(z_3|z_2) p(z_2|z_1) p(z_1)$$

- Notice that:  $z_4 \perp\!\!\!\perp \{z_1, z_2\} | z_3$
- We want to make inference on  $z_4$

$$p(z_4) = \sum_{z_1, z_2, z_3} p(z_1, z_2, z_3, z_4)$$

# Variable Elimination

## simple PGM



- A trivial solution is to go over all possible combinations of values!

$$p(z_4) = \sum_{z_1} \sum_{z_2} \sum_{z_3} p(z_1, z_2, z_3, z_4)$$

- Generally, if each of the  $m$  variables has  $k$  possible values, we'd have a complexity of  $O(k^{m-1})$
- This quickly becomes intractable (just remember our *trivial* example with a mixture model)  $\rightarrow$  NP-hard problem



# Variable Elimination

## simple PGM

- We can take advantage of the PGM structure

$$\begin{aligned} p(z_4) &= \sum_{z_1, z_2, z_3} p(z_1, z_2, z_3, z_4) \\ &= \sum_{z_1, z_2, z_3} p(z_4|z_3) p(z_3|z_2) p(z_2|z_1) p(z_1) \\ &= \sum_{z_3} p(z_4|z_3) \sum_{z_2} p(z_3|z_2) \sum_{z_1} p(z_2|z_1) p(z_1) \end{aligned}$$

- In a chain graph, the complexity reduces to  $O(mk^2)$

# Variable Elimination

simple PGM



$$p(z_4) = \sum_{z_3} p(z_4|z_3) \sum_{z_2} p(z_3|z_2) \sum_{z_1} p(z_2|z_1) p(z_1)$$

# Variable Elimination

## simple PGM




$$\begin{aligned} p(z_4) &= \sum_{z_3} p(z_4|z_3) \sum_{z_2} p(z_3|z_2) \sum_{z_1} p(z_2|z_1) p(z_1) \\ &= \sum_{z_3} p(z_4|z_3) \sum_{z_2} p(z_3|z_2) f_a(z_2) \end{aligned}$$

- Notice that  $f_a(z_2)$  is a function of  $z_2$  (also called a *factor*). For example, a CPT with the probabilities of the  $k$  values of  $z_2$
- We just “got rid” of  $z_1$  by marginalizing over its values
- Same as in last lecture, when we marginalized over  $z$  for implementing LDA in STAN...

# Variable Elimination

## simple PGM



$$\begin{aligned} p(z_4) &= \sum_{z_3} p(z_4|z_3) \sum_{z_2} p(z_3|z_2) \sum_{z_1} p(z_2|z_1) p(z_1) \\ &= \sum_{z_3} p(z_4|z_3) \sum_{z_2} p(z_3|z_2) f_a(z_2) \\ &= \sum_{z_3} p(z_4|z_3) f_b(z_3) \end{aligned}$$

# Variable Elimination

## simple PGM



$$\begin{aligned} p(z_4) &= \sum_{z_3} p(z_4|z_3) \sum_{z_2} p(z_3|z_2) \sum_{z_1} p(z_2|z_1) p(z_1) \\ &= \sum_{z_3} p(z_4|z_3) \sum_{z_2} p(z_3|z_2) f_a(z_2) \\ &= \sum_{z_3} p(z_4|z_3) f_b(z_3) \\ &= f_c(z_4) \end{aligned}$$

## Variable Elimination (VE)


- Time complexity is exponential in size of largest factor
  - Each  $f_i$  is a factor
  - Size of factor is number of variables it depends on
- Order is vital (and sometimes a complicated problem)!
- Observed variables

$$p(z|x = k) = \frac{p(z, x = k)}{p(x = k)}$$

- Perform VE on  $p(z, x = k)$  (i.e. we fix  $x=k$ ) and then on  $p(x = k) = \sum_z p(z, x = k)$
- **Only acyclic graphs!**

# Playtime!

- Apply the Variable Elimination algorithm to the following graph
- We want to infer  $p(d|b = 1)$

 $p(a|b,c)$ 

	$a = 0$	$a = 1$
$b = 0, c = 0$	0.7	0.3
$b = 0, c = 1$	0.3	0.7
$b = 1, c = 0$	0.5	0.5
$b = 1, c = 1$	0.1	0.9

 $p(c)$ 

$c = 0$	$c = 1$
0.7	0.3

 $p(b)$ 

$b = 0$	$b = 1$
0.4	0.6

 $p(d|a)$ 

	$d = 0$	$d = 1$
$a = 0$	0.6	0.4
$a = 1$	0.2	0.8

# Belief propagation in Polytrees

Pearl's algorithm

- A graph is a polytree if (and only if) there is at most one simple path between any two nodes,  $v_i$  and  $v_k$




Figure: Polytree




Figure: Not polytree




Figure: ?




Figure: ?




# Belief propagation in Polytrees

Pearl's algorithm

- We want to solve  $p(\mathbf{z}|\mathbf{x})$ , i.e. the conditional probability of all variables  $\mathbf{z} = \{z_1, z_2, \dots, z_V\}$ , given evidence  $\mathbf{x} = \{x_1, x_2, \dots, x_E\}$
- Graph structure allows for incremental inference steps
  - Like in Variable Elimination...
- Some nodes are clearly specified from beginning
  - Evidence  $x_i = k$
  - Priors
- The other nodes can build on them
- Information flows as **messages** between nodes

# Belief propagation in Polytrees


Pearl's algorithm



- $\pi$  messages - from parent to child

# Belief propagation in Polytrees

Pearl's algorithm




- $\pi$  messages - from parent to children
- Represent the current belief about the parent - **causal evidence**
- After receiving all parents' messages, one can know more of the child

$$\pi(z_j) = \sum_{z_i \in \text{parent}(z_j)} p(z_j | z_1, z_2, \dots) \prod_{z_i \in \text{parent}(z_j)} \pi_{z_i, z_j}(z_i)$$

# Belief propagation in Polytrees

Pearl's algorithm



- **$\lambda$  messages** - from child to parent

$$p(z_j | z_i)$$

- Represents the belief about the parent's value - **diagnostic evidence**


# Belief propagation in Polytrees

Pearl's algorithm

- **$\lambda$  messages** - from child to parent

$$p(z_i|z_j)$$

- Represents the belief about the parent's value - **diagnostic evidence**
- After receiving all children's messages, one can know more of the parent




$$\lambda(z_j) = \prod_{z_i \in \text{child}(z_j)} \lambda_{z_i, z_j}(z_j)$$

# Belief propagation in Polytrees

Pearl's algorithm


- But we still need the formulas for the messages!

$$\pi_{z_i, z_j}(z_i) = \pi(z_i) \prod_{v \neq j} \lambda_{z_v, z_i}(z_i)$$



- $\lambda$  message when we have  $p(z_i|z_j)$  (only one parent,  $z_j$ )

$$\lambda_{z_i, z_j}(z_j) = \sum_{z_i} \lambda(z_i) p(z_i|z_j)$$




# Belief propagation in Polytrees

Pearl's algorithm

- Generic formula for  $\lambda$  messages:

$$\lambda_{z_i, z_j}(z_j) = \sum_{z_i} \lambda(z_i) \sum_{z_v \in \mathbf{z} \setminus z_j} p(z_i | z_1, z_2, \dots, z_j) \prod_{v \neq j} \pi_{z_v, z_i}(z_v)$$



## Pearl's algorithm - Initialization step

- For all  $x_i \in \mathbf{x}$  (evidence nodes):
  - $\lambda(x_i) = 1$ , wherever  $x_i = e_i$ , 0 otherwise (i.e. change the CPD tables, or dirac delta function)
  - $\pi(x_i) = 1$ , wherever  $x_i = e_i$ , 0 otherwise (i.e. change the CPD tables, or dirac delta function)
- For all nodes,  $z_i$ , without parents:
  - $\pi(z_i) = p(z_i)$ , the prior
- For all nodes,  $z_i$ , without children:
  - $\lambda(z_i) = 1$




# Pearl's algorithm

- Iterate until no change occurs
  - 1 For each node  $z_i$ , if it has received all  $\pi_{*,z_i}$  messages, calculate  $\pi(z_i)$
  - 2 For each node  $z_i$ , if it has received all  $\lambda_{*,z_i}$  messages, calculate  $\lambda(z_i)$
  - 3 For each node  $z_i$ , if  $\pi(z_i)$  is calculated, and all  $\lambda_{*,z_i}$  messages have been received (except from  $z_j$ ), calculate  $\pi_{z_i,z_j}(z_i)$  and send the message to  $z_j$
  - 4 For each node  $z_i$ , if  $\lambda(z_i)$  is calculated, and all  $\pi_{*,z_i}$  messages have been received (except from  $z_j$ ), calculate  $\lambda_{z_i,z_j}(z_j)$  and send the message to  $z_j$
- Compute  $\text{Bel}(z_i) \propto \lambda(z_i) \pi(z_i)$  and normalize, for all desired nodes

## Example

- Let's apply Belief Propagation to the following graph
- We want to infer  $p(d|b = 1)$



$p(a|b,c)$

	$a = 0$	$a = 1$
$b = 0, c = 0$	0.7	0.3
$b = 0, c = 1$	0.3	0.7
$b = 1, c = 0$	0.5	0.5
$b = 1, c = 1$	0.1	0.9

$p(c)$

$c = 0$	$c = 1$
0.7	0.3

$p(b)$

$b = 0$	$b = 1$
0.4	0.6

$p(d|a)$

	$d = 0$	$d = 1$
$a = 0$	0.6	0.4
$a = 1$	0.2	0.8

## Example

- From the initialization, we have

$\pi(c) = p(c)$	
$c = 0$	$c = 1$
0.7	0.3

$\pi(b)$	
$b = 0$	$b = 1$
0	1


  

$\lambda(b)$	
$b = 0$	$b = 1$
0	1

$\lambda(d)$	
$b = 0$	$b = 1$
1	1

## Example

- We start with the evidence and prior



$$\pi_{c,a}(c) = \pi(c) \prod_{x \neq a} \lambda_{x,c}(c) = \pi(c)$$


$\pi_{c,a}(c)$	
$c = 0$	$c = 1$
0.7	0.3

$$\pi_{b,a}(b) = \pi(b)$$

$\pi_{b,a}(b)$	
$b = 0$	$b = 1$
0	1

## Example

- Now,  $\pi(a)$




$$\pi(a) = \sum_{b,c} p(a|b,c) \pi_{b,a}(b) \pi_{c,a}(c)$$

For  $a=1$ :

$$\begin{aligned}\pi(a=1) &= \sum_{b=\{0,1\}} \sum_{c=\{0,1\}} p(a=1|b,c) \pi_{b,a}(b) \pi_{c,a}(c) \\ &= \sum_{c=\{0,1\}} p(a=1|b=1,c) \pi_{c,a}(c) \\ &= p(a=1|b=1,c=0) \pi_{c,a}(c=0) + p(a=1|b=1,c=1) \pi_{c,a}(c=1) \\ &= 0.5 \times 0.7 + 0.9 \times 0.3 = 0.62, \text{ so } \pi(a=1) = 0.62 \text{ and } \pi(a=0) = 0.38\end{aligned}$$

## Example

- Finally, we reach  $d$



Since there are no more parents, we have  $\pi_{a,d}(a) = \pi(a)$

$$\pi_{a,d}(a)$$

$a = 0$	$a = 1$
0.38	0.62

## Example

$$\pi(d) = \sum_{a=\{0,1\}} p(d|a) \pi_{a,d}(a)$$

$p(d a)$		
	$d = 0$	$d = 1$
$a = 0$	0.6	0.4
$a = 1$	0.2	0.8

$\pi_{a,d}(a)$	
$a = 0$	$a = 1$
0.38	0.62

- Trivially becomes:


$\pi(d)$	
$b = 0$	$b = 1$
0.352	0.648

- $\text{Bel}(d) = \pi(d)\lambda(d)$
- In this case, the solution is trivially:

$$p(d|b=1) = \pi(d)$$

- But why stop here? We can propagate back, and also get  $p(a|b=1)$  and  $p(c|b=1)$

## Example



$$\lambda_{d,a}(a) = \sum_{d=\{0,1\}} \lambda(d)p(d|a)$$

- Since  $\lambda(d) = 1$ , we have:


$$\lambda_{d,a}(a) = \sum_{d=\{0,1\}} p(d|a) = 1$$

- Notice that we have a unnormalized table!

$\lambda_{d,a}(a)$	
$a = 0$	$a = 1$
1	1



## Example



- Since  $\lambda(a) = 1$ , we have:

$$\lambda(c) = \lambda_{a,c}(c) = \sum_{a=\{0,1\}} \lambda(a) p(a|b=1, c) = \sum_{a=\{0,1\}} p(a|b=1, c) = 1$$


- Therefore:

- $\text{Bel}(c) = \pi(c) \lambda(c) = \pi(c) = p(c)$
- $\text{Bel}(a) = \pi(a) \lambda(a) = \pi(a) = \sum_c p(a|b=1, c) \pi_{c,a}(c)$ , as calculated before

- Makes sense, right (notice the independence of the graph!)?

## Example

- But, what if there's a new evidence,  $d = 0$ ?



- We only need to update  $d$  and propagate back!
- A new initialization gives:

$\lambda(d)$	
$d = 0$	$d = 1$
1	0

$\lambda(a)$	
$a = 0$	$a = 1$
0.6	0.2

$$\lambda_{d,a}(a) = \sum_{d=\{0,1\}} \lambda(d)p(d|a) = p(d=0|a)$$

$$\lambda(a) = \lambda_{d,a}(a)$$

$$\begin{aligned} \lambda_{a,c}(c) &= \sum_{a,b=\{0,1\}} \lambda(a)p(a|b,c)\pi_{b,a}(b) = \\ &= \sum_{a=\{0,1\}} \lambda(a)p(a|b=1,c) \end{aligned}$$

$$\lambda(c) = \lambda_{a,c}(c)$$

## Example

- So,
- $\text{Bel}(a) = \alpha \lambda(a) \pi(a)$ , with  $\alpha$  the normalizing factor

$\pi(a)$	
$a = 0$	$a = 1$
0.38	0.62

$\lambda(a)$	
$a = 0$	$a = 1$
0.6	0.2

- $\alpha = (0.38 * 0.6 + 0.62 * 0.2)^{-1} = 2.84$

$$p(a|b = 1, d = 0) = \text{Bel}(a)$$

$a = 0$	$a = 1$
0.65	0.35



$$p(c|b = 1, d = 0) = \text{Bel}(c)$$

$c = 0$	$c = 1$
0.8	0.2

# Belief Propagation

- A *forward-backward* pass of the BP algorithm gives us the exact inference for every variable
- Updating can be done incrementally
- Plates can be treated by expansion (given that it is still a polytree)
- We used discrete variables, but the reasoning is the same for continuous
  - Summations become integrals
  - Need to derive a new function for each step (easy in some cases, particularly with exponential family)
- In theory, it is still NP-Hard, but in practice, it is often linear time
- All of the above is valid for polytrees

## The problem of non-polytrees







- Creates an (infinite) cycle
  - Clique Trees
  - Loopy belief propagation

# Clique trees

- We aggregate nodes together (maximal cliques), and create a new graph, which is a polytree
- We may not get individual variable marginals, but instead joint distributions of subsets of variables
  - ...but you can marginalize individually in the end!
- One can create a clique tree just by following the Variable Elimination algorithm:
  - Each factor  $f_i$  corresponds to a node in the tree
  - When factor  $f_i$  uses the result of factor  $f_j$ , then there is an undirected edge between the two.
  - It becomes a polytree

# Clique trees

- Algorithm for creating a polytree
  - Given a set of factors, construct the undirected graph  $H_\phi$
  - Triangulate  $H_\phi$ , to construct a chordal graph<sup>2</sup>  $H^*$
  - Find cliques in  $H^*$ , and make each one a node in a cluster graph
  - Run the maximum spanning tree algorithm on the cluster graph to construct a tree




<sup>2</sup>In a chordal graph, all cycles of four or more vertices have a chord, which is an edge that is not part of the cycle but connects two vertices of the cycle.

# Loopy Belief Propagation

- Apply BP on the original graph in the following way:
  - ① Initialize all messages to 1
  - ② Run BP algorithm as before (start with evidence/priors)
  - ③ Each node sends messages in parallel (i.e. when it sends to one child/parent, it sends to all children/parents)
  - ④ Loop until convergence
- Works very well when there are some loops, but not a fully connected graph
- With one complete loop, the MAP should be correct
- If  $p(\mathbf{z})$  is jointly Gaussian, Loopy Belief Propagation will converge to the correct marginals



# Playtime!



$p(c)$

$c = 0$	$c = 1$
0.7	0.3

$p(b|a)$

	$b = 0$	$b = 1$
$a = 0$	0.3	0.7
$a = 1$	0	1

$p(a|c,d)$

	$a = 0$	$a = 1$
$c = 0, d = 0$	0.5	0.5
$c = 0, d = 1$	0.9	0.1
$c = 1, d = 0$	0.1	0.9
$c = 1, d = 1$	0	1

$p(d)$

$d = 0$	$d = 1$
0.8	0.2

$p(e|a)$

	$e = 0$	$e = 1$
$a = 0$	0.2	0.8
$a = 1$	0.7	0.3

- Consider the graph and CPTs above
- Perform forward-backward Belief Propagation, assuming the evidence that  $d = 1, e = 1$

## Some final notes

- There are actually many tools that do BP
  - Just check: <https://www.cs.ubc.ca/~murphyk/Software/bnsoft.html>
- BP is often applied on the tractable sub-parts of your model, for example combining belief from different sub-models:
  - A Random Forest, RF provides  $p_{\text{RF}}(y)$ , and a Logistic Regression model, LR, provides  $p_{\text{LR}}(y)$ :

$$\text{Bel}(y) = \alpha p_{\text{RF}}(y) p_{\text{LR}}(y)$$

- Also called *Bayesian Model Averaging*
- If you have gentle analytical forms for the sub-models (e.g. normal distribution for outputs), the combination is straightforward
- You can even treat sub-models as “priors” that provide to your random variables, and combine many in a complex model!

## References

- Kautz and Lowd, CSE 573: Introduction to Artificial Intelligence.  
[https://courses.cs.washington.edu/courses/cse573/05au/10\\_26\\_pearl.ppt](https://courses.cs.washington.edu/courses/cse573/05au/10_26_pearl.ppt)
- Koller and Friedman, (2009) Koller, D., and Friedman, N. Probabilistic graphical models: principles and techniques. MIT press. (2009) (Chapters 9 and 10).
- Christopher, M. Bishop. Pattern Recognition and Machine Learning. Springer-Verlag New York, 2016.